

Rodin

Approach

- learn 3D avatar generation off multi-view renderings from 'Blender Synthetic Pipeline':
Computer software that gives 2D images of 3D objects from known viewpoints

3D Training Data Representation

We need a way to represent 3D neural radiance fields that:

- Can input to Neural Networks
- Is memory efficient & fast

Vanilla NeRF approach would take too long to create a synthetic dataset

Tri-Plane Representation

- Each 3D point $p \in \mathbb{R}^3$ is mapped to 3 planar coordinates p_{uv} , p_{uw} , p_{vw} for planes uv , uw , vw representing each axis.
- The density & colour of each point is encoded on a planar level as embeddings y_{uv} , y_{uw} , y_{vw}
- The overall embedding is computed as $y_p = y_{uv} p_{uv} + y_{uw} p_{uw} + y_{vw} p_{vw}$ which represents the colour & density of point p from different viewing angles as a number vector
- A lightweight MLP decoder then can extract the colour & density from each y_p given p and viewing direction d
- We also apply a positional encoding here to shift focus to high frequency details.

3D Diffusion Model

- Learns tri-plane features (tri-plane encoding of the colour & density) $P(y)$

where $y = (y_{uv}, y_{uw}, y_{vw})$ based off diffusion process



- Initial Approach: we have 3 planes $\in \mathbb{R}^{W \times H \times C}$, concatenate the 3 planes to make 1 image $\in \mathbb{R}^{W \times H \times 3C}$ & use current SOTA 2D diffusion models

Problem: we lose spatial properties since each plane is no longer separate but on top of each other

3D Aware Concatenation

- For each point y , for the plane y_{uv} , add in some information that tells it about the other planes at y : y_{uw}, y_{vw} . Now do for the other planes.
 - by taking the line intersection of the 2 planes, embedding it & then concatenating this info into y_{uv} .
- Allows colour & density at y_{uv} to depend on y_{uw} & y_{vw}

Latent Conditioning

- Use CLIP to represent the front view of the image as a latent vector (which CLIP can understand with text)
- Now as the diffusion model constructs images, it checks with CLIP that it is creating the correct coherent features. CLIP can now also manipulate the generative model with text input.

→ AdaGN framework

Model Upsampler

- After diffusion model outputs a 64×64 image, we train a network to upsample to 256×256 images to enhance quality by feeding the network diffusion images & ground truth.
 - Also train a convolution resizer to find high-frequency details that NeRF misses taking the images to 1024×1024 .
- all done using tri-plane coordinate system

What can the model do?

Avatar Portrait

- Input: Single input image
output: 3D renderable avatar

} CLIP just feeds the image's own latent vector to diffusion network

Text → Avatar

- Input: text description
output: 3D renderable avatar

Customisable Avatar

→ CLIP takes current latent vector z & adds relevant features s to output $z+s$

- Input: input image / current avatar & text description of changes (ie. "add glasses")
- Output: 3D renderable avatar